Text Analysis in R

Introduction
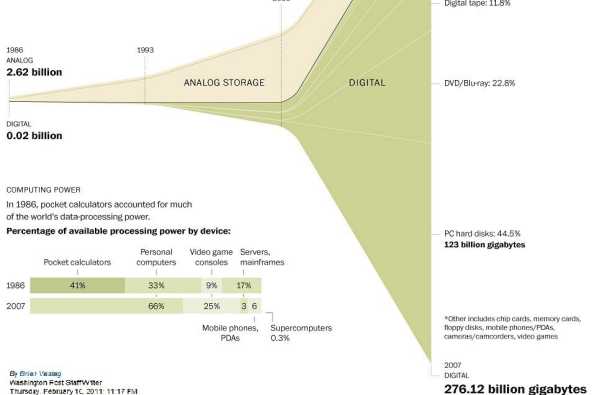
Helge Liebert

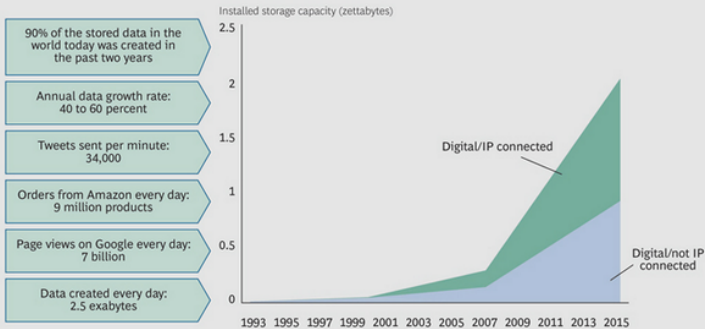THE WORLD'S CAPACITY TO STORE INFORMATION

This chart shows the world's growth in storage capacity for both analog data (books, newspapers, videotapes, etc.) and digital (CDs, DVDs, computer hard drives, smartphone drives, etc.)

In gigabytes or estimated equivalent

1986
ANALOG
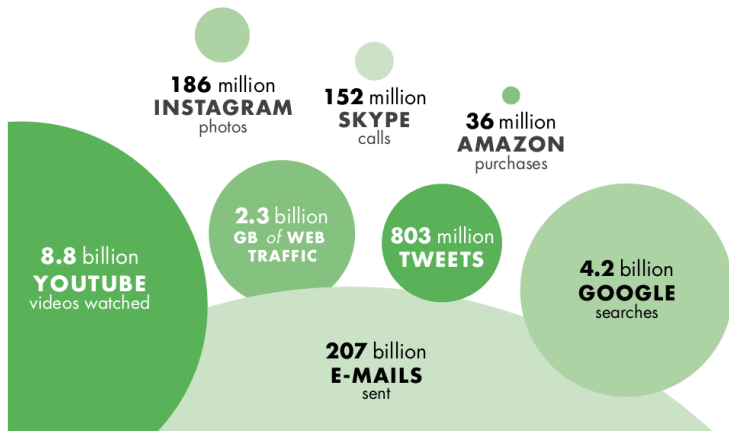**2.62 billion**

1993

2000

ANALOG STORAGE

DIGITAL

DIGITAL
0.02 billion

2007
ANALOG
18.86 billion gigabytes

Paper, film, audiotape and vinyl: 6.2%
Analog videotapes: 93.8%

ANALOG
DIGITAL

Other digital media: 0.8%*
Portable media players, flash drives: 2%
Portable hard disks: 2.4%

CDs and minidisks: 6.8%

Computer servers and
mainframe hard disks: 8.9%

Digital tape: 11.8%

DVD/Blu-ray: 22.8%

PC hard disks: 44.5%
123 billion gigabytes

*Other includes chip cards, memory cards,
floppy disks, mobile phones/PDAs,
cameras/camcorders, video games

2007
DIGITAL
276.12 billion gigabytes

COMPUTING POWER

In 1986, pocket calculators accounted for much
of the world's data-processing power.

Percentage of available processing power by device:

	Pocket calculators	Personal computers	Video game consoles	Servers, mainframes	Mobile phones, PDAs	Supercomputers
1986	41%	33%	9%	17%		
2007		66%	25%	3	6	0.3%

EXHIBIT 2 | In 2015, More Than Half of All Data Will Have an IP Address

Installed storage capacity (zettabytes)

- 90% of the stored data in the world today was created in the past two years
- Annual data growth rate: 40 to 60 percent
- Tweets sent per minute: 34,000
- Orders from Amazon every day: 9 million products
- Page views on Google every day: 7 billion
- Data created every day: 2.5 exabytes

Digital/IP connected

Digital/not IP connected

1993 1995 1997 1999 2001 2003 2005 2007 2009 2011 2013 2015

Sources: Martin Hilbert and Priscilla Lopez, "The World's Technological Capacity to Store, Communicate, and Compute Information," *Science*, February 2011; BCG estimates.

b. A typical day in the life of the internet

186 million
INSTAGRAM
photos

152 million
SKYPE
calls

36 million
AMAZON
purchases

2.3 billion
GB *of* **WEB TRAFFIC**

803 million
TWEETS

8.8 billion
YOUTUBE
videos watched

4.2 billion
GOOGLE
searches

207 billion
E-MAILS
sent

Sources: World Development Indicators (World Bank, various years); WDR 2016 team; http://www.internetlivestats.com/one-second/ (as compiled on April 4, 2015). Data at http://bit.do/WDR2016-FigO_4.

Note: In panel a, for some years data for electricity are interpolated from available data. GB = gigabytes.

Introduction

- 90% of data today has been created in the last two years.
- 235 million emails sent per day.
- 3.3 million Facebook posts created every minute.
- 3.8 million Google searches performed each minute.
- 1.7 megabytes of new information created every second, per person.

Introduction

- 90% of data today has been created in the last two years.
- 235 million emails sent per day.
- 3.3 million Facebook posts created every minute.
- 3.8 million Google searches performed each minute.
- 1.7 megabytes of new information created every second, per person.

➥ An immense amount of data, new and old, is recorded as text.
➥ More generally, much of this data is unstructured.

Structured vs. unstructured

Structured data

- Adheres to a defined data model.
- Examples: Tables, spreadsheets, relational databases, …

Unstructured data

- Does *not* adhere to a defined data model.
- Typically text-heavy.
- Examples: Text feeds, speech transcripts, audio, images …

Structured vs. unstructured

Structured data

- Adheres to a defined data model.
- Examples: Tables, spreadsheets, relational databases, …

Semi-structured data

- Does *not* adhere to a formal data model,
- … *but* contains tags or semantic mark-up.
- Examples: JSON, XML, emails, tagged text, …

Unstructured data

- Does *not* adhere to a defined data model.
- Typically text-heavy.
- Examples: Text feeds, speech transcripts, audio, images …

Text as data

- Text differs from other, traditional forms of data.
- Text is inherently *unstructured* and *high-dimensional*.
- One of the major fields of application of machine learning methods.
- Fast-growing field. Many new techniques developed in industry.
- Recent applications in economics and other social sciences.

This lecture

This lecture covers techniques for unstructured data.

- Methods for wrangling data.
- ➡ When unstructured $\approx$ dirty (or differently structured).

This lecture

This lecture covers techniques for unstructured data.

- Methods for wrangling data.
- ➥ When unstructured ≈ dirty (or differently structured).

- Methods for analyzing data which are naturally unstructured.
- ➥ No rectangular (or graph) structure, no well-defined relations between data elements.

Focus points

Focus on three main points.

1. Processing and transforming un-/semi-structured data.
2. Representing inherently unstructured text data.
3. Analyzing text data and using models to discover structure. (Supervised and unsupervised learning.)

Outline

Dates

Wednesday	16.02.2022	08.30-16.30	A230
Thursday	17.02.2022	08.30-16.30	A230
Friday	18.02.2022	08.30-16.30	A230

Schedule

08.30-10.00	Lecture
10.00-10.30	Break
10.30-12.00	Lecture
12.00-13.30	Lunch
13.30-15.00	Lecture
15.00-15.30	Break
15.30-16.30	Lecture

Technical requirements

- All class material is available online:
  https://github.com/hliebert/course-text-analysis-in-r.

- The lab materials can be accessed online:
  Jupyter notebooks
  Rstudio server

- Feel free to run the lab material locally on your own computer. Clone or download the course repository to get started.

- It will run on Windows, Mac or Linux (if the dependencies are satisfied).

Programs

Minimal

- A browser.

Local

- R.
- Editor or GUI (RStudio, VScode with R plugin, Jupyter, Emacs+ESS, ...).
- Run the R install script provided with the class material to install the R package dependencies and the R Kernel for Jupyter notebooks.

Optional

- Jupyter notebooks. Install Anaconda (or its smaller miniconda version). On Linux, you can also use python pip.

Assignment

1. Independent project (100%)

- Deadline: 15.03.2022.
- More details during the course of the lecture.

Primary references

- The course covers relatively broad and diverse topics, no single reference. Seminal references in the slides.
- Primary and secondary references below.
- Hastie et al. and Jurafsky & Martin books are available online (use newest 3rd edition draft of J & M).

📄 Gentzkow, M., B. Kelly, and M. Taddy (2019). Text as Data. *Journal of Economic Literature* 57(3), 535–574. DOI: `10/gf7rd5`.

📕 Hastie, T., R. Tibshirani, and J. Friedman (2001). *The Elements of Statistical Learning: Data Mining, Inference, and Prediction.* Ed. by R. Tibshirani and J. H. (H. Friedman. New York.

📕 Jurafsky, D. and J. H. Martin (2009). *Speech and Language Processing: An Introduction to Natural Language Processing, Computational Linguistics, and Speech Recognition.* 2nd ed. Prentice Hall Series in Artificial Intelligence. Upper Saddle River, N.J: Pearson Prentice Hall.

Secondary references

- Reference material, further reading, applied or introductory text books.

📓 Baumer, B., D. Kaplan, and N. Horton (2017). *Modern Data Science with R.* CRC.

📕 Casella, G. and R. L. Berger (2001). *Statistical Inference.* Second. Duxbury Press.

📓 Chacon, S. and B. Straub (2014). *Pro Git.* Apress.

📓 Goldberg, Y. (2017). *Neural Network Methods for Natural Language Processing.* Synthesis Lectures on Human Language Technologies 37. San Rafael: Morgan & Claypool Publishers.

📕 James, G., D. Witten, T. Hastie, and R. Tibshirani (2015). *An Introduction to Statistical Learning with Applications in R.* Springer.

📓 Matloff, N. (2011). *The Art of R Programming: A Tour of Statistical Software Design.* No Starch Press.

📓 Shotts, W. E. (2019). *The Linux Command Line: A Complete Introduction.* Second edition. San Francisco: No Starch Press.

Secondary references

📕 Silge, J. and D. Robinson (2017). *Text Mining with R: A Tidy Approach*. First edition. Boston: O'Reilly.

📕 Wasserman, L. (2006). *All of Nonparametric Statistics*. Springer.

📕 Wasserman, L. (2010). *All of Statistics: A Concise Course in Statistical Inference*. Springer Texts in Statistics. New York, NY: Springer.