

Text Analysis in R

Introduction

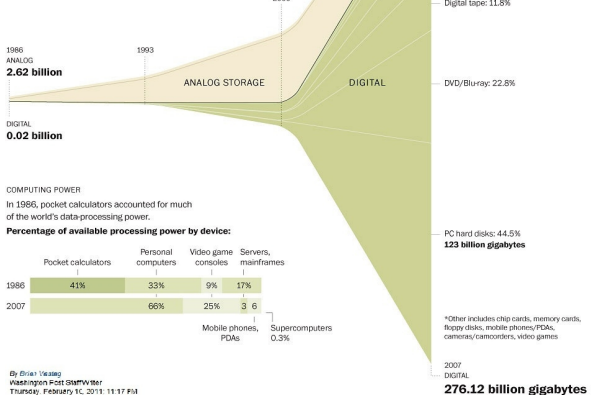
Helge Liebert

Worldwide data storage capacity

THE WORLD'S CAPACITY TO STORE INFORMATION


This chart shows the world's growth in storage capacity for both analog data (books, newspapers, videotapes, etc.) and digital (CDs, DVDs, computer hard drives, smartphone drives, etc.)

In gigabytes or estimated equivalent



Data, then and now

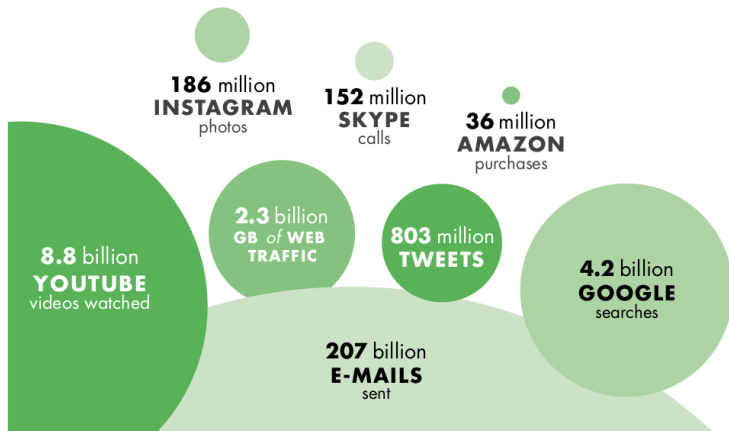
EXHIBIT 2 | In 2015, More Than Half of All Data Will Have an IP Address



Sources: Martin Hilbert and Priscilla Lopez, "The World's Technological Capacity to Store, Communicate, and Compute Information," *Science*, February 2011; BCG estimates.

Data, then and now

b. A typical day in the life of the internet



Sources: World Development Indicators (World Bank, various years); WDR 2016 team; <http://www.internetlivestats.com/one-second/> (as compiled on April 4, 2015). Data at http://bit.do/WDR2016-FigO_4.

Note: In panel a, for some years data for electricity are interpolated from available data. GB = gigabytes.

Introduction

- 90% of data today has been created in the last two years.
- 235 million emails sent per day.
- 3.3 million Facebook posts created every minute.
- 3.8 million Google searches performed each minute.
- 1.7 megabytes of new information created every second, per person.

Introduction

- 90% of data today has been created in the last two years.
 - 235 million emails sent per day.
 - 3.3 million Facebook posts created every minute.
 - 3.8 million Google searches performed each minute.
 - 1.7 megabytes of new information created every second, per person.
- ➡ An immense amount of data, new and old, is recorded as **text**.
- ➡ More generally, much of this data is **unstructured**.

Structured vs. unstructured

Structured data

- Adheres to a defined data model.
- Examples: Tables, spreadsheets, relational databases, ...

Unstructured data

- Does *not* adhere to a defined data model.
- Typically text-heavy.
- Examples: Text feeds, speech transcripts, audio, images ...

Structured vs. unstructured

Structured data

- Adheres to a defined data model.
- Examples: Tables, spreadsheets, relational databases, ...

Semi-structured data

- Does *not* adhere to a formal data model,
- ... *but* contains tags or semantic mark-up.
- Examples: JSON, XML, emails, tagged text, ...

Unstructured data

- Does *not* adhere to a defined data model.
- Typically text-heavy.
- Examples: Text feeds, speech transcripts, audio, images ...

- Text differs from other, traditional forms of data.
- Text is inherently *unstructured* and *high-dimensional*.
- One of the major fields of application of machine learning methods.
- Fast-growing field. Many new techniques developed in industry.
- Recent applications in economics and other social sciences.

This lecture

This lecture covers techniques for unstructured data.

- Methods for wrangling data.
- ➡ When unstructured \approx dirty (or differently structured).

This lecture

This lecture covers techniques for unstructured data.

- Methods for wrangling data.
 - ➡ When unstructured \approx dirty (or differently structured).
- Methods for analyzing data which are naturally unstructured.
 - ➡ No rectangular (or graph) structure, no well-defined relations between data elements.

Focus points

Focus on three main points.

1. Processing and transforming un-/semi-structured data.
2. Representing inherently unstructured text data.
3. Analyzing text data and using models to discover structure.
(Supervised and unsupervised learning.)

Outline

1. Introduction

Representation

2. Regular expressions and pattern matching

3. Representing text as data

n-gram modeling approaches

4. Supervised models for text data

5. Unsupervised models for text data

Information retrieval and distributional language models

6. Distributional models of meaning

7. Vector space representations

Assignment

Dates

Wednesday	15.02.2022	08.30-16.30	C130
Thursday	16.02.2022	08.30-16.30	C130
Friday	17.02.2022	08.30-16.30	C130

Schedule

08.30-10.00	Lecture
10.00-10.30	Break
10.30-12.00	Lecture
12.00-13.30	Lunch
13.30-15.00	Lecture
15.00-15.30	Break
15.30-16.30	Lecture

Technical requirements

- All class material is available online:
<https://github.com/hliebert/course-text-analysis-in-r>.
- The lab materials can be accessed online:
[Jupyter notebooks](#)
[Rstudio server](#)
- Feel free to run the lab material locally on your own computer. Clone or download the course repository to get started.
- It will run on Windows, Mac or Linux (if the dependencies are satisfied).

Programs

Minimal

- A browser.

Local

- R.
- Editor or GUI (RStudio, VScode with R plugin, Jupyter, Emacs+ESS, ...).
- Run the R install script provided with the class material to install the R package dependencies and the R Kernel for Jupyter notebooks.

Optional

- Jupyter notebooks. Install Anaconda (or its smaller miniconda version). On Linux, you can also use python pip.

Assignment

1. Independent project (100%)

- Deadline: 31.03.2023.
- More details during the course of the lecture.

Primary references

- The course covers relatively broad and diverse topics, no single reference. Seminal references in the slides.
- Textbooks are available online.
- Use latest 3rd edition draft of Jurafsky & Martin.
- For general ML/Stats, consult Hastie & Tibshirani or Murphy.



Gentzkow, M., B. Kelly, and M. Taddy (2019). Text as Data. *Journal of Economic Literature* 57(3), 535–574. DOI: [10/gf7rd5](https://doi.org/10/gf7rd5).



Hastie, T., R. Tibshirani, and J. Friedman (2001). *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. Ed. by R. Tibshirani and J. H. (H. Friedman. New York.



Jurafsky, D. and J. H. Martin (2009). *Speech and Language Processing: An Introduction to Natural Language Processing, Computational Linguistics, and Speech Recognition*. 2nd ed. Prentice Hall Series in Artificial Intelligence. Upper Saddle River, N.J: Pearson Prentice Hall.



Murphy, K. P. (2022). *Probabilistic Machine Learning: An introduction*. MIT Press. URL: probml.ai.

Secondary references

- Reference material, further reading, applied or introductory text books.



Baumer, B., D. Kaplan, and N. Horton (2017). *Modern Data Science with R*. CRC.



Casella, G. and R. L. Berger (2001). *Statistical Inference*. Second. Duxbury Press.



Chacon, S. and B. Straub (2014). *Pro Git*. Apress.



Goldberg, Y. (2017). *Neural Network Methods for Natural Language Processing*. Synthesis Lectures on Human Language Technologies 37. San Rafael: Morgan & Claypool Publishers.



James, G., D. Witten, T. Hastie, and R. Tibshirani (2015). *An Introduction to Statistical Learning with Applications in R*. Springer.



Matloff, N. (2011). *The Art of R Programming: A Tour of Statistical Software Design*. No Starch Press.



Shotts, W. E. (2019). *The Linux Command Line: A Complete Introduction*. Second edition. San Francisco: No Starch Press.

Secondary references



Silge, J. and D. Robinson (2017). *Text Mining with R: A Tidy Approach*. First edition. Boston: O'Reilly.



Wasserman, L. (2006). *All of Nonparametric Statistics*. Springer.



Wasserman, L. (2010). *All of Statistics: A Concise Course in Statistical Inference*. Springer Texts in Statistics. New York, NY: Springer.